POSTDOC STATEMENT | Naman Gupta

Through this postdoc fellowship, I want to challenge the coloniality and epistemic erasure embedded in AI and work towards making AI more value-sensitive, diverse and safe for everyone.

I will create a framework that will evaluate epistemic and physical violence witnessed by people involved in AIs' development and operations. I will contrast how Big Tech AI models and infrastructure disproportionately impacts people and land in and outside the imperial core (generally, WEIRD populations) and the periphery (generally, the majority world) and those who sit at their boundaries (Indigenous peoples in Western countries). This framework aims to unearth the extractive logics embedded in every layer of BigTech AI lifecycle, from data collection to model training, from extraction of minerals that form the physical AI infrastructure to privacy policies of the AI interfaces used by consumers. The framework will provide opportunities that shape AI safety that incorporates sustainability and sovereignty.

Scholars have noted that academic research reflect an imperial impulse [1–3]. Consequently, AI technologies center Western hegemonic understandings of safety, ethics, responsibility, and privacy and perpetuates epistemic violence [4] and enforces universalization and monoculture [5–7]. However, without engaging with theories of change that center decoloniality, sovereignty, and liberation from the Majority World [8–10], AI ethics and safety research may result in neoliberal siloed interventions that fail to address the ethics and privacy violations of peoples harmed by the enormous power of the Big Tech AI enterprise. For example, trust and safety teams in Big Tech companies often tout their products as private, ethical, carbon-neutral, whereas, the same companies actively exploit miners who extract rare earth minerals in Congo and assemble their products in India. BigTech companies create their own rules for trust, safety, and ethics and show meaningless numbers that they abide by it. On a similar vein, what does privacy mean for exploited factory workers in India who watch abusive content or wear cameras to train AI robots that will soon replace them? Similarly, ChatGPT promises privacy settings that ensure users' data is not used for model training, yet also, legally allows intimate partner abusers to generate fake nude images of their victims without consent.

Drawing on decolonial and feminist literature [10, 11], my goal is to make these contestations of BigTech AI products visible so that it can inform AI safety and ethical frameworks beyond positivist notions of safety metrics and benchmarks. A narrow definition of ethics, privacy, and safety only brings surface level interventions that rarely addresses the underlying structural violence and prevents communities from achieving long-term healing and liberation.

Methods. My computing and critical theory expertise allows me to view technologies not just as a black box, but as a combination of abstract social and technical components. I study impacts of violence through personal technologies (e.g., mobile accounts, AirTags) [NDSS-LBW 2024, CHI-LBW 2024, USENIX 2024, CCS 2025, In-review USENIX 2027, CSCW 2025, CHI 2026, In-review CSCW 2027] and from physical infrastructure (e.g., AI data centers, AMBER alerts) [In-review CSCW 2027]. I will surface the underlying neoliberalism and coloniality in AI safety research across security and privacy, HCI, ethics, psychology, and critical computing spaces. To do this, I will build on my workshop that calls for extending the definition of violence that transcends disciplinary boundaries [In-review DIS 2026] and counter-computing [In-review CHI 2027] resistance movements that counter the hegemonic technologies and infrastructures. I will generate keywords of concepts that capture the power differentials among people involved in the development of an AI model using the 'Calculating Empires' framework [23]. Then, I will use automated scraping methods [In-review CSCW 2027] and snowball sampling to collect a wide collection of relevant research articles, news reports, and, podcasts that formally or informally define ethics and privacy in global contexts. I will categorize these definitions into mode of harms, the actors involved (including people and the land), their impacts on actors, and the mitigation and harm-reduction strategies to negotiate and address ethical and privacy concerns. These will include communities strategies to resist, refuse, reclaim, and build new infrastructure for care, sovereignty, and liberation. Next, to ensure diversity in the definitions and by situating them in place and space, I will interview key stakeholders embedded in a variety of aspects of the AI lifecycle. For example, organizers fighting AI data centers, survivors of gender-based violence who use AI for support, immigrants who face social media surveillance, or, researchers that work on AI safety research. The interview will utilize the Meta-Relationality protocol [24] to surface a diverse set of orientations around affect (feelings about different kinds of AI?), mitigation strategy (what theory of change informs their response to AI?), time (how does the time of shape their stance?), and agency (what kind of leverage is possible and required?). I utilized this framework to conduct a workshop with top CS and Math researchers [HLF 2026]. I sensed rising tensions and uncertainty among scholars in the midst of planetary crises and scientific acceleration and corporatization and lack of perceived agency to shape the future of AI.

Findings and Impact. Findings from the study will result in a composite set of diverse contestations that show the agreements and contradictions across value-systems, social positions, and proximity to AI knowledge creation and usage. These contestations are necessary to inform the privacy interventions so that they center agency and liberation from AI-mediated violence: whether that may be refusal on ethical grounds, regulatory action or guardrails, resistance movement, decentralized alternatives, sabotage from within academia or corporations, or, eco-social symbiotic technologies. This work reimagines how sovereign and culturally-sensitive, sustainable, sovereign, and ethical AI can manifest and benefit the communities.

My Publications

[NDSS-LBW 2024] **Naman Gupta**, Rose Ceccio, Majed Almansoori, and Rahul Chatterjee. "Analyzing the Patterns and Behavior of Users When Detecting and Preventing Tech-enabled Stalking". In: 2023. URL: https://doi.org/10/gr8rnx.

[CHI-LBW 2024] **Naman Gupta**, Sanchari Das, Kate Walsh, and Rahul Chatterjee. "A Critical Analysis of the Prevalence of Technology-Facilitated Abuse in US College Students". In: *Extended Abstracts of the 2024 CHI Conference on Human Factors in Computing Systems*. 2024. URL: https://doi.org/10.1145/3613905.3652036.

[USENIX 2024] **Naman Gupta**, Kate Walsh, Sanchari Das, and Rahul Chatterjee. ""I really just leaned on my community for support": Barriers, Challenges, and Coping Mechanisms Used by Survivors of {Technology-Facilitated} Abuse to Seek Social Support". In: 2024. URL: https://dl.acm.org/doi/10.5555/3698900.3699179.

[CCS 2025] Julia Nonnenkamp, **Naman Gupta**, Abhimanyu Dev Gupta, and Rahul Chatterjee. "Hidden in Plain Bytes: Investigating Interpersonal Account Compromise with Data Exports". In: *2025 ACM SIGSAC Conference on Computer and Communications Security*. 2025. URL: https://doi.org/10.1145/3719027.3765147.

[In-review USENIX 2027] Julia Nonnenkamp, **Naman Gupta**, Angela Wang, Sophie Stephenson, and Rahul Chatterjee. "LEStrADE: Investigating Account Compromise using Data Exports for Survivors of Technology-Facilitated Abuse". In: *review at USENIX Security*. 2026. URL: https://naman.github.io/assets/files/papers/lestrade.pdf.

[CSCW 2025] Sophie Stephenson, **Naman Gupta**, Akhil Polamarasetty, Kyle Huang, David Youssef, Kayleigh Cowan, and Rahul Chatterjee. "Legal Evidence of Technology-Facilitated Abuse in Wisconsin: Surfacing Barriers Within and Beyond the Courtroom." In: *Computer Supported Cooperative Work and Social Computing* 🏆 **Honorable Mention (top 5%)**. 2025. URL: https://doi.org/10.1145/3757622.

[CHI 2026] Sophie Stephenson, **Naman Gupta**, Kyle Huang, David Youssef, Kayleigh Cowan, and Rahul Chatterjee. "Trauma-Informed Digital Evidence Collection:A Design Inquiry into Evidence Practices for Technology-Facilitated Abuse in Intimate Partner Violence". In: *2026 CHI Conference on Human Factors in Computing Systems* 🏆 **Best Paper (top 1%)**. 2026. URL: https://doi.org/10.1145/3772318.3790296.

[In-review CSCW 2027] **Naman Gupta**, Connie W. Chau, Sophie Stephenson, Rahul Chatterjee, and Kate Walsh. "Lessons from Trauma-Informed Training on Technology-Facilitated Abuse for Gender-Based Violence Advocates". In: *review at CSCW* (2026). URL: https://arxiv.org/abs/2610.02688.

[In-review CSCW 2027] **Naman Gupta**, Sophie Stephenson, Chung Chi Yeung, Wei Ting Wu, Jeneile Luebke, Kate Walsh, and Rahul Chatterjee. ""Lighting The Way For Those Not Here": How Technology Researchers Can Help Fight the Missing and Murdered Indigenous Relatives (MMIR) Crisis". In: *review at CSCW 2027*. 2025. URL: https://doi.org/10.48550/arXiv.2601.17966.

[In-review DIS 2026] Jasmine C Foriest, Alicia DeVrio, Li Qiwei, Nimra Ahmed, **Naman Gupta**, Jensine Raihan, Franchesca Spektor, and Justine Zhang. "Violence as a Boundary Object: Language, Power, and Design Justice". In: *Review at ACM DIS Designing Interactive Systems*. 2026.

[In-review CHI 2027] **Naman Gupta**, Freddy Reiber, Jasmine C. Foriest, Ailsa Robertson, Jordan Taylor, Palak Dudani, Ghadeer Awwad, Shreya Chowdhary, David Gamba, Eugene Kim, and Katie Seaborn. "Counter-Computing: Digital Resistances and Global Solidarities". In: *Review at ACM CHI Conference on Human Factors in Computing Systems*. 2027.

[HLF 2026] **Naman Gupta**. "Polycrisis in the Age of AI and Scientific Acceleration". In: 2026.

References

[1] Paul Dourish and Scott D. Mainwaring. "Ubicomp's colonial impulse". In: *Proceedings of the 2012 ACM Conference on Ubiquitous Computing*. UbiComp '12. Association for Computing Machinery, 2012.

[2] Syed Mustafa Ali. "A Brief Introduction to Decolonial Computing". In: *XRDS* (2016).

[3] Jensine Raihan. "The Colonial Legacy of the Human in Human-Computer Interaction: Implications for Critical Design Approaches". In: *Proc. Sixth Decenn. Aarhus Conf. Comput. X Crisis*. AAR '25. Association for Computing Machinery, 2025.

[4] Paola Ricaurte. *Artificial Intelligence and the Feminist Decolonial Imagination*. en. GSCC: 0000015 2026-08-21T19:09:04.150Z 0 Read_Status: New Read_Status_Date: 2026-08-21T19:09:03.367Z. 2022. URL: https://botpopuli.net/?post_type=post&p=5927 (visited on 08/21/2026).

[5] Seth Lazar and Alondra Nelson. "AI safety on whose terms?" In: *Science* 381.6654 (July 2023). 11 citations (Crossref) [2024-06-30] tex.ids= lazarAISafetyWhose2023a, pp. 138–138. DOI: 10/gsgsmv. URL: https://www.science.org/doi/10.1126/science.adi8982 (visited on 07/15/2023).

[6] Boaventura de Sousa Santos. *Epistemologies of the South: Justice Against Epistemicide*.

[7] Lisa Messeri and M. J. Crockett. "Artificial intelligence and illusions of understanding in scientific research". en. In: *Nature* 627.8002 (Mar. 2024). Read_Status: New Read_Status_Date: 2026-11-08T23:12:35.236Z, pp. 49–58. ISSN: 1476-4687. DOI: 10.1038/s41586-024-07146-0. URL: https://www.nature.com/articles/s41586-024-07146-0 (visited on 11/08/2026).

[8] Arturo Escobar. *Designs for the Pluriverse: Radical Interdependence, Autonomy, and the Making of Worlds*. Duke University Press, 2018.

[9] Morgan Ndlovu. "Decoloniality Is Agency". In: (2025).

[10] Paola Ricaurte. "Ethics for the Majority World: AI and the Question of Violence at Scale". In: *Media Cult. Soc.* (2022).

[11] Warmhold Jan Thomas Mollema. "Decolonial AI as Disenclosure". en. In: *Open Journal of Social Sciences* 12.2 (Feb. 2024). 0 citations (Crossref/DOI) [2024-08-17] Number: 2, pp. 574–603. DOI: 10/gt7bjt. URL: https://www.scirp.org/journal/paperinformation?paperid=131479 (visited on 08/17/2024).

[23] Kate Crawford and Vladan Joler. *Calculating Empires: A Genealogy of Technology and Power since 1500*. 2023.

[24] Meta-Relationality Institute and Gesturing Towards Decolonial Futures Collective. *Clearing the Field: A Relational Protocol for Difficult Conversations about AI*. 2025.